

Content from Deliverable_1.ipynb

```
# MSCS 634 - Project Deliverable 1
# Advanced Data Mining for Data-Driven Insights and Predictive Modeling
# Dataset: Stroke Prediction Dataset (Kaggle)

# Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Set display and style options
pd.set_option('display.max_columns', None)
sns.set(style='whitegrid')

# Load Dataset
df = pd.read_csv('/content/healthcare-dataset-stroke-data.csv')
```

```
# Inspect structure
print("Dataset shape:", df.shape)
df.head()

# Basic Info
df.info()
df.describe(include='all')

# Check missing values
print("\nMissing values per column:\n", df.isnull().sum())

# Data Cleaning
# Drop 'id' column (not useful for modeling)
df.drop(columns=['id'], inplace=True)

# Handle missing values in 'bmi' (replace with median)
df['bmi'].fillna(df['bmi'].median(), inplace=True)
```

```
Dataset shape: (5110, 12)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
 0   id          5110 non-null   int64  
 1   gender      5110 non-null   object  
 2   age         5110 non-null   float64 
 3   hypertension 5110 non-null   int64  
 4   heart_disease 5110 non-null   int64  
 5   ever_married 5110 non-null   object  
 6   work_type    5110 non-null   object
```

```

7   Residence_type    5110 non-null  object
8   avg_glucose_level 5110 non-null  float64
9   bmi               4909 non-null  float64
10  smoking_status    5110 non-null  object
11  stroke            5110 non-null  int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB

Missing values per column:
id          0
gender      0
age         0
hypertension 0
heart_disease 0
ever_married 0
work_type    0
Residence_type 0
avg_glucose_level 0
bmi        201
smoking_status 0
stroke      0
dtype: int64
/tmp/ipython-input-2214330086.py:17: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

```

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the ope

```
df['bmi'].fillna(df['bmi'].median(), inplace=True)
```

```

# Check for duplicates
duplicates = df.duplicated().sum()
print("Duplicate records:", duplicates)
df.drop_duplicates(inplace=True)

# Fix categorical inconsistencies (e.g., work_type typo)
df['work_type'].replace({'Govt_jov': 'Govt_job'}, inplace=True)

# Confirm cleaning
df.info()

```

```

Duplicate records: 0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender          5110 non-null  object 
 1   age             5110 non-null  float64
 2   hypertension    5110 non-null  int64  
 3   heart_disease   5110 non-null  int64  
 4   ever_married    5110 non-null  object 
 5   work_type       5110 non-null  object 
 6   Residence_type  5110 non-null  object 
 7   avg_glucose_level 5110 non-null  float64
 8   bmi             5110 non-null  float64
 9   smoking_status  5110 non-null  object 
 10  stroke          5110 non-null  int64  
dtypes: float64(3), int64(3), object(5)
memory usage: 439.3+ KB

```

```
/tmp/ipython-input-1039141086.py:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.  
  
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the ope
```

```
df['work_type'].replace({'Govt_jov': 'Govt_job'}, inplace=True)
```


```
# Exploratory Data Analysis (EDA)  
# Distribution of Target Variable  
plt.figure(figsize=(6,4))  
sns.countplot(x='stroke', data=df, palette='pastel')  
plt.title('Stroke Occurrence Distribution')  
plt.show()  
  
# Numeric Distributions  
num_cols = ['age', 'avg_glucose_level', 'bmi']  
df[num_cols].hist(bins=20, figsize=(10,6), color='lightblue')  
plt.suptitle('Distribution of Numerical Features')  
plt.show()  
  
# Correlation Heatmap  
plt.figure(figsize=(8,6))  
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='Blues')  
plt.title('Correlation Heatmap')  
plt.show()  
  
# Relationship between Age, BMI, and Stroke  
plt.figure(figsize=(7,5))  
sns.boxplot(x='stroke', y='age', data=df, palette='coolwarm')  
plt.title('Age vs Stroke')  
plt.show()  
  
plt.figure(figsize=(7,5))  
sns.boxplot(x='stroke', y='bmi', data=df, palette='coolwarm')  
plt.title('BMI vs Stroke')  
plt.show()  
  
# Categorical Breakdown  
plt.figure(figsize=(8,4))  
sns.countplot(x='gender', hue='stroke', data=df)  
plt.title('Gender vs Stroke')  
plt.show()  
  
plt.figure(figsize=(8,4))  
sns.countplot(x='work_type', hue='stroke', data=df)  
plt.title('Work Type vs Stroke')  
plt.xticks(rotation=45)  
plt.show()
```






```
/tmp/ipython-input-1133343763.py:4: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```


```
sns.countplot(x='stroke', data=df, palette='pastel')
```




Distribution of Numerical Features



Correlation Heatmap








/tmp/ipython-input-1133343763.py:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='stroke', y='age', data=df, palette='coolwarm')
```



```
0 1  
stroke  
/tmp/ipython-input-1133343763.py:27: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.  
sns.boxplot(x='stroke', y='bmi', data=df, palette='coolwarm')
```



Deliverable 2 content

2000
1

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split, KFold, cross_validate
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score, mean_squared_error

data = load_diabetes()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = pd.Series(data.target)

numeric_features = X.columns.tolist()

preprocessor = Pipeline([
    ("poly", PolynomialFeatures(degree=2, include_bias=False)),
    ("scaler", StandardScaler())
])

models = {
    "LinearRegression": Pipeline([("prep", preprocessor), ("reg", LinearRegression())]),
    "Ridge": Pipeline([("prep", preprocessor), ("reg", Ridge())]),
    "Lasso": Pipeline([("prep", preprocessor), ("reg", Lasso(max_iter=10000))])
}

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

cv = KFold(n_splits=5, shuffle=True, random_state=42)
results = {}

for name, model in models.items():
    cv_res = cross_validate(model, X_train, y_train, scoring=["r2", "neg_mean_squared_error"], cv=cv)
    mse = -cv_res["test_neg_mean_squared_error"]
    rmse = np.sqrt(mse)
    results[name] = {
        "R2_CV": cv_res["test_r2"].mean(),
        "RMSE_CV": rmse.mean()
    }

print(pd.DataFrame(results).T)

test_results = {}
trained_models = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    trained_models[name] = model
    pred = model.predict(X_test)
    test_results[name] = {
        "R2 Test": r2_score(y_test, pred),
        "RMSE Test": np.sqrt(mean_squared_error(y_test, pred))
    }
```

```
        "RMSE_Test": np.sqrt(mean_squared_error(y_test, pred))
    }

test_df = pd.DataFrame(test_results).T
print(test_df)

best_model_name = test_df["R2_Test"].idxmax()
best_model = trained_models[best_model_name]

pred = best_model.predict(X_test)

plt.figure(figsize=(6,6))
plt.scatter(y_test, pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], "--")
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted")
plt.show()

residuals = y_test - pred
plt.figure(figsize=(8,4))
sns.histplot(residuals, kde=True)
plt.title("Residuals")
plt.show()

plt.figure(figsize=(8,4))
plt.scatter(pred, residuals, alpha=0.6)
plt.axhline(0, linestyle="--", color="black")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residuals vs Predicted")
plt.show()
```


	R2_CV	RMSE_CV
Linear Regression	0.38059	62.504792
Ridge	0.382059	60.465540
Lasso	0.445447	57.133522

```
import warnings, sys, os

# Hide Python warnings
warnings.filterwarnings("ignore")

# Hide ALL stderr warnings from libraries like jupyter_client
sys.stderr = open(os.devnull, "w")
import logging
logging.disable(logging.CRITICAL)
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    confusion_matrix, ConfusionMatrixDisplay, roc_curve, auc,
    accuracy_score, f1_score, classification_report
)
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, DBSCAN

# For association rules
from mlxtend.frequent_patterns import apriori, association_rules

# reproducibility
RANDOM_STATE = 42

print('libraries imported')
```


libraries imported

```
#Load dataset & exploratory data analysis
data = load_breast_cancer(as_frame=True)
df = data.frame.copy()

df['target'] = data.target

df.head()
# Cell: quick EDA
print(df.shape)
print(df['target'].value_counts())
df.describe().T.head()
# Cell: correlation heatmap (top 20 features by variance)
plt.figure(figsize=(12,8))
```

```
cols = df.columns[:-1] # exclude target
corr = df[cols].corr()
# plot only a subset for readability
sns.heatmap(corr.iloc[:20,:20], cmap='RdBu_r', center=0)
plt.title('Feature correlation (subset)')
plt.show()
```



```
# Preprocessing
X = df.drop(columns=['target'])
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print('train/test shapes:', X_train.shape, X_test.shape)
```

train/test shapes: (455, 30) (114, 30)

```
# train decision tree
dt = DecisionTreeClassifier(random_state=RANDOM_STATE)
dt.fit(X_train_scaled, y_train)

y_pred_dt = dt.predict(X_test_scaled)

# metrics
acc_dt = accuracy_score(y_test, y_pred_dt)
f1_dt = f1_score(y_test, y_pred_dt)
print('Decision Tree - Accuracy:', acc_dt, 'F1:', f1_dt)

# display confusion matrix
ConfusionMatrixDisplay.from_estimator(dt, X_test_scaled, y_test)
plt.title('Decision Tree - Confusion Matrix')
plt.show()

# ROC curve
y_proba_dt = dt.predict_proba(X_test_scaled)[:,1]
fpr_dt, tpr_dt, _ = roc_curve(y_test, y_proba_dt)
roc_auc_dt = auc(fpr_dt, tpr_dt)
plt.plot(fpr_dt, tpr_dt, label=f'DT (AUC = {roc_auc_dt:.3f})')


# Cell: k-NN
knn = KNeighborsClassifier()
knn.fit(X_train_scaled, y_train)
y_pred_knn = knn.predict(X_test_scaled)
acc_knn = accuracy_score(y_test, y_pred_knn)
f1_knn = f1_score(y_test, y_pred_knn)
print('k-NN - Accuracy:', acc_knn, 'F1:', f1_knn)

ConfusionMatrixDisplay.from_estimator(knn, X_test_scaled, y_test)
plt.title('k-NN - Confusion Matrix')
plt.show()


# ROC for k-NN
y_proba_knn = knn.predict_proba(X_test_scaled)[:,1]
fpr_knn, tpr_knn, _ = roc_curve(y_test, y_proba_knn)
roc_auc_knn = auc(fpr_knn, tpr_knn)
plt.plot(fpr_knn, tpr_knn, label=f'k-NN (AUC = {roc_auc_knn:.3f})')
```

```
# finish ROC plot
plt.plot([0,1],[0,1], '--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
plt.show()
```


Decision Tree - Accuracy: 0.9122807017543859 F1: 0.9285714285714286



k-NN - Accuracy: 0.956140350877193 F1: 0.9655172413793104



```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV

param_dist = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 5, 10, 20],
    'max_features': ['sqrt', 'log2'],
}

rf = RandomForestClassifier(random_state=RANDOM_STATE)

rs = RandomizedSearchCV(
    rf,
    param_distributions=param_dist,
    n_iter=10,
    scoring='f1',
    cv=5,
    random_state=42,
```